

[Get started](#)[Open in app](#)[Follow](#)

537K Followers



You have **1** free member-only story left this month. [Sign up for Medium and get an extra one](#)

How To Plot Time Series



Nadim Kawwa · Nov 3, 2019 · 5 min read ★



Photo by [Chris Lawton](#) on [Unsplash](#)

Dealing with time series can be one of the most insightful parts of exploratory data analysis, if done right. In This post, we are going to use the checkin log from the Yelp Dataset to explore trends across different time periods using Pandas and Matplotlib.

Data Acquisition

After downloading the data, we need to know what to use. The block below shows a sample entry from the `checkin.json` file based on the Yelp Documentation:

```
{
  // string, 22 character business id, maps to business in
  business.json
  "business_id": "tnhfdv5I18EaGSXZGiuQGg"

  // string which is a comma-separated list of timestamps for each
  checkin, each with format YYYY-MM-DD HH:MM:SS
  "date": "2016-04-26 19:49:16, 2016-08-30 18:36:57, 2016-10-15
02:45:18, 2016-11-18 01:54:50, 2017-04-20 18:39:06, 2017-05-03
17:58:02"
}
```

We can read the input file with pandas read_json method with arguments

`orient=columns` and `Lines=True`.

Upon reading the data, our dataframe looks something like this:

	business_id	date
0	--1UhMGODdWsrMastO9DZw	2016-04-26 19:49:16, 2016-08-30 18:36:57, 2016...
1	--6MefnULPED_I942VcFNA	2011-06-04 18:22:23, 2011-07-23 23:51:33, 2012...
2	--7zmmkVg-IMGaXbuVd0SQ	2014-12-29 19:25:50, 2015-01-17 01:49:14, 2015...
3	--8LPVSo5i0Oo61X01sV9A	2016-07-08 16:43:30
4	--9QQLMTbFzLJ_oT-ON3Xw	2010-06-26 17:39:07, 2010-08-01 20:06:21, 2010...

The `date` column entries are strings such that each date is separated by a comma. By looking at them we can tell that the format is indeed `YYYY-MM-DD HH:MM:SS`.

Data Wrangling

For readability, we want our dataframe to look more like this:

date	Business_id
2011-06-04 18:22:23	123aAdf
2011-07-23 23:51:33	--6MefnULPED_I942VcFNA

The unfolded event log makes more sense from an EDA perspective whereas the format in the JSON format makes more sense for memory storage.

We can obtain this result by creating a dictionary where the keys correspond to dates and the values correspond to business_id. Here we assume that the dates are granular enough so that no two dates are the same. Keep in mind that the [python documentation](#) does not allow duplicate keys within a single dictionary.

The implementation can be done as follows:

```

1  #create dict that maps ID
2  bus_to_check = {}
3
4  for bus, dates in zip(df['business_id'], df['date']):
5      dates_as_list = dates.split(',')
6      bus_to_check[bus] = pd.to_datetime(dates_as_list)
7
8  #create the dict with dates as keys
9  check_to_bus = dict( (v,k) for k in bus_to_check for v in bus_to_check[k])
10
11 #create dataframe from dict
12 event_df = pd.DataFrame.from_dict(check_to_bus, orient='index', columns=['business_id'])
13 #sort by the index for neatness
14 event_df = event_df.sort_index(ascending=True)

```

df_to_eventlog_yelp.py hosted with ❤ by GitHub

[view raw](#)

Seasonality


Say we want to know what are the total checkins for all the years available. We will use this as a gateway to introduce the pandas Grouper which can be used inside the groupby method. For more information about frequency aliases refer to the pandas docs.

```
1 #sum over years
2 checkin_year_sigma = event_df.groupby(pd.Grouper(freq='Y')).count()
3 #plot year v total checkins
4 plt.plot(checkin_year_sigma.index, checkin_year.business_id)
```

yelp_groupby_year_sigma.py hosted with ❤ by GitHub

[view raw](#)

The code returns this neat plot:



it is a good start, however what if we want to dive deeper? We will zoom in on the year 2014, though any other year will do. Selecting date ranges is an easy one-liner:

```
1 #get 2014-2015 date indices
2 df_2014 = checkin_day[(checkin_day.index>'2014-01-01') & (checkin_day.index<'2015-01-01')]
```

yelp_select_year_2014.py hosted with ❤ by GitHub

[view raw](#)

A first indicator of seasonality is to look at the weekend, since the data is from the USA & Canada the weekend is Saturday and Sunday. Bear in mind that the weekend varies by country, for example Dubai's weekend is Friday and Saturday.

We therefore want to detect the days that are a weekend and write a customizable function for it:

```

1  def find_weekend_indices(datetime_array, weekend=5):
2      """
3      Returns all indices of Saturdays & Sundays in a datetime array
4      datetime_array(pandas) = pandas datetime array
5      weekend(int) = assume weekend starts at day=5=Saturday
6      """
7      #empty list to store indices
8      indices = []
9
10
11     for i in range(len(datetime_array)):
12         # get day of the week with Monday=0, Saturday=5, Sunday=6
13         if datetime_array[i].weekday() >= weekend:
14             indices.append(i)
15
16     return indices

```

find_weekend_indices.py hosted with ❤ by GitHub

[view raw](#)

The function above returns indices that are weekends. Next for plotting we define another helper function to highlight spans of the plot that would correspond to these indices with the help of pyplot axvspan:

```

1  def highlight_datetimes(indices, ax, df, facecolor='green', alpha_span=0.2):
2      """
3      Highlights all weekends in an axes object
4      indices(list) = list of Saturdays and Sundays indices corresponding to dataframe
5      ax(matplotlib) = pyplot object
6      df(pandas) = pandas dataframe
7      """
8      i = 0
9      #iterate over indices
10     while i < len(indices)-1:
11         #highlight from i to i+1
12         ax.axvspan(df.index[indices[i]],
13                   df.index[indices[i] + 1],
14                   facecolor=facecolor,
15                   edgecolor='none',
16                   alpha=alpha_span)

```

17 `i += 1`

highlight_datetimes.py hosted with ❤ by GitHub

[view raw](#)

Finally we can use these two functions inside a larger one that takes in a dataframe:

```
1  def plot_datetime(df, title="Yelp Businesses Checkins",
2                      highlight=True,
3                      weekend=5,
4                      ylabel="Number of Visits",
5                      saveloc=None,
6                      facecolor='green',
7                      alpha_span=0.2):
8      """
9      Draw a plot of a dataframe that has datetime object as its index
10     df(pandas) = pandas dataframe with datetime as indeces
11     highlight(bool) = to highlight or not
12     title(string) = title of plot
13     saveloc(string) = where to save file
14     """
15
16     #instantiate fig and ax object
17     fig, axes = plt.subplots(nrows=1, ncols=1, sharex=True,figsize=(15,5))
18
19     #draw all columns of dataframe
20     for v in df.columns.tolist():
21         axes.plot(df[v], label=v, alpha=.8)
22
23     if highlight:
24         #find weekend indeces
25         weekend_indices = find_weekend_indices(df.index, weekend=5)
26         #highlight weekends
27         highlight_datetimes(weekend_indices, axes, df, facecolor)
28
29     #set title and y label
30     axes.set_title(title, fontsize=12)
31     axes.set_ylabel(ylabel)
32     axes.legend()
33     plt.tight_layout()
34
35     #add xaxis gridlines
36     axes.xaxis.grid(b=True, which='major', color='black', linestyle='--', alpha=1)
37
38     #savefig if
```

```


39     if saveloc:
40         fig.savefig(saveloc)
41     plt.show()

```


plot_datetime.py hosted with ❤ by GitHub

[view raw](#)

The result is this neat plot:



The plot above matches the intuition that people go out more on weekends than weekdays. To take a better look at that buildup we can take a closer look at April 2014 using the same function.



What if we want to show a more granular view of April? Since each entry in the original dataframe is a single checkin, we can groupby 30 minute frequencies and count as follows:

```

1  #get checkins by 30min intervals
2  checkin_halfhour = event_df.groupby(pd.Grouper(freq='30min')).count()

```

```


3 #select interval
4 checkin_halfhour_midapril = checkin_halfhour[(checkin_halfhour.index>'2014-04-02 00:00:00')
5                                               & (checkin_halfhour.index<'2014-04-06 00:00:00')]
6 #plot
7 plot_help.plot_datetime(checkin_halfhour_midapril,
8                           title='Yelp Businesses Checkins Mid April 2014',
9                           saveloc='plots/checkin_2014_april_30min.png')

```

yelp_midapril_30min.py hosted with ❤ by GitHub

[view raw](#)

The result is the plot below:



Plotting Confidence Intervals

We saw how we can zoom in on time periods to get a sense of seasonality. However, businesses also want to know what is the expected number of checkins for a given day and how it might vary. For example, we might be solving for a regression.

Therefore we need to build plots that show the average checkin for a given day of the week and some sort of confidence interval.

We will use the `checkin_halfhour` defined in the previous section to extract the day of the week:

```

1 # add a column including time only
2 checkin_halfhour['Time']=checkin_halfhour.index.time
3
4 #column for name of day of week
5 checkin_halfhour['DayOfWeek'] = checkin_halfhour.index.weekday_name

```

checkin_halfhour_dayofweek.py hosted with ❤ by GitHub

[view raw](#)

Our dataframe now looks as follows:

	business_id	Time	DayOfWeek
2010-01-15 22:30:00	1	22:30:00	Friday
2010-01-15 23:00:00	1	23:00:00	Friday
2010-01-15 23:30:00	1	23:30:00	Friday
2010-01-16 00:00:00	0	00:00:00	Saturday
2010-01-16 00:30:00	1	00:30:00	Saturday

Next, we build the `day_avg` dataframe using pandas aggregate method. For our purposes we will aggregate:

- Sum
- Mean
- Standard Deviation

The code block below does this for us:

```

1  #group by day of week and aggregate by sum+mean+stdf
2  day_avg=checkin_halfhour.groupby(['DayOfWeek','Time']).agg([np.sum, np.mean, np.std])
3
4  #drop level in columns
5  day_avg.columns = day_avg.columns.droplevel()
6
7  #extract upper and lower bound as mean +/- 2STD
8  day_avg['upper'] = day_avg['mean'] + 2*day_avg['std']
9  day_avg['lower'] = day_avg['mean'] - 2*day_avg['std']
10
11 #extract day of week name
12 day_avg['day']=day_avg.reset_index()['DayOfWeek'].tolist()
```

yelp_day_avg.py hosted with ❤ by GitHub

[view raw](#)

We now have a dataframe we can use to draw intervals:

		sum	mean	std	upper	lower	day
DayOfWeek	Time						
Friday	00:00:00	92875	201.902174	71.791131	345.484435	58.319913	Friday
	00:30:00	99487	216.276087	78.435046	373.146180	59.405994	Friday
	01:00:00	102484	222.791304	78.231877	379.255057	66.327551	Friday
	01:30:00	102034	221.813043	78.483890	378.780823	64.845264	Friday
	02:00:00	96047	208.797826	73.036509	354.870843	62.724809	Friday

How we define the interval depends on the distribution of the data and the business value of narrow/larger intervals. In our case, we want to generate the expected checkin for each day of the week and grab all the information within two standard deviations.

Doing so is a matter of watching out for the levels in the index of `day_avg` as shown below:




```


1  #days of week
2  days = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']
3
4  for day in days:
5      ax=day_avg[day_avg['day']==day].unstack(level=0)['mean'].plot(figsize=(10,5))
6      day_avg[day_avg['day']==day].unstack(level=0)['upper'].plot(ax=ax)
7      day_avg[day_avg['day']==day].unstack(level=0)['lower'].plot(ax=ax)
8      ax.set_title('Yelp Checkin Trend for {}'.format(day))
9      ax.legend(['Mean', '+ 2 STD', '- 2 STD'])
10     plt.tight_layout()
11     plt.savefig('plots/checkin_trend_'+day+'.png')
```

yelp_checkin_expected_dayofweek.py hosted with ❤ by GitHub

[view raw](#)

The resulting plots reveal a lot in terms of checkin behavior, for illustration we show Monday, Wednesday, and Saturday.





Conclusion

The amount and depth of time series analysis depends on the problem we are trying to solve for. For example, we might want to evaluate restaurants, or businesses that only open from 8AM to 3PM such as cafés.

An objective of EDA is to prepare for feature extraction. With the exercise above we can start thinking about features such as the most common checkin time, or checkins by day.

Sign up for The Daily Pick

By Towards Data Science

Hands-on real-world examples, research, tutorials, and cutting-edge techniques delivered Monday to Thursday. Make learning your daily ritual. [Take a look](#)

Your email

Get this newsletter

By signing up, you will create a Medium account if you don't already have one. Review our [Privacy Policy](#) for more information about our privacy practices.

Data Science Python Education Business Intelligence Yelp

[About](#) [Help](#) [Legal](#)